

XRONOS: an open data infrastructure for archaeological chronology

Joe Roe 

University of Bern

joseph.a.roe@unibe.ch

Clemens Schmid 

Max Planck Institute for Geoanthropology

Setareh Ebrahimiabareghi 

University of Bern

Caroline Heitz 

University of Bern

Martin Hinz 

University of Bern

ABSTRACT XRONOS (<https://xronos.ch>) is an open data infrastructure for the backbone of the archaeological record – chronology. It aims to provide open access to published radiocarbon dates and other chronometric data from any period, anywhere in the world. Collating a number of recent regional and global compilations of dates, XRONOS offers the most comprehensive radiocarbon database yet published, with over 350,000 date and 75,000 site records. It also provides a foundation for expanding the systematic collection of chronometric information beyond radiocarbon, with support for typological and dendrochronological dates and a generalisable data model that can be adapted to other methods of absolute dating. Automated and semi-automated quality control processes ensure that data from diverse sources is continuously integrated and standardised, making it significantly easier to filter data of interest and reducing the need for manual data cleaning by end users. In this paper we describe the concept and implementation of XRONOS in relation to the state of the art in chronometric data-sharing, and evaluate its potential as a general-purpose open repository and curation platform for archaeological chronology.

KEYWORDS open data; chronology; chronometry; radiocarbon dating; dendrochronology; typological dating

Introduction

- Chronology is the backbone of the archaeological record – we need it to be open!

Although ‘open data’ is often thought to be a new phenomenon in archaeology (CITE), it has a long history in radiometric dating. Starting in the 19XXs, radiocarbon laboratories shared and compiled as ‘date lists’ published on a regular basis

*Corresponding author.

in the journal *Radiocarbon* (CITE). However, as the number of labs and volume of radiocarbon dates being produced grew, this paper-based format became impractical: the last ‘date list’ appeared in 19XX, and unfortunately was not replaced by another form of systematic data-sharing or dissemination. Additionally, because the date lists were sourced from radiocarbon laboratories directly—not those who collected the sample—they typically included only very limited contextual information. Thus, even though radiocarbon data comes from a relatively limited number of sources (**IntCalLabList**) and has relatively standardised reporting conventions (Millard 2014), in practice the only way to produce aggregated datasets recently has been to manually search through relevant literature for dates reported by the submitter of the sample. This already laborious process is further hampered by a significant inconsistency in how much authors adhere to measurement reporting conventions (Millard 2014), the lack of any convention on the reporting of contextual information, and weak or nonexistent disciplinary norms regarding the responsibility to publish results openly in a timely fashion.

- Move to open compilations in the last ~10 years
- Range of other chronological information used by archaeologists is largely untouched
- From uploading CSVs to (social) infrastructure

State of the art

Compilations of radiocarbon dates

- Where do radiocarbon dates come from and how are they published?
- History of systematic compilations
- Summary of existing databases
 - FIGURE: graph of number of dates/databases through time
 - FIGURE: map of radiocarbon database coverage
 - TABLE: radiocarbon databases
- Problems with current state of the art

Open data infrastructures

The profusion of radiocarbon compilations over the last decade has naturally prompted many to think more globally:

- The first available radiocarbon ‘metadatabase’ was **c14bazAAR** (CITE), an R package that provides an index of openly published radiocarbon databases and a common interface for retrieving them and performing basic data cleaning. Because c14bazAAR downloads data from its original source repositories, rather than mirroring it, it only includes resources that have been published in a fully open and machine-accessible format. Despite this limitation, it has global coverage and is still the largest collection of radiocarbon dates available (see FIG<empty citation>).
- Another indexical approach is the **IntChron** project (CITE), which exposes data from multiple sources and exposes them with a common JSON-based web interface. The IntChron specification is open, meaning that radiocarbon labs or compilation projects can implement it independently and thereby allow end users to access their data through a common interface (though to our knowledge it has so far only been adopted by databases associated with the

Oxford Radiocarbon Lab). The JSON format also lends itself to the implementation of wrapper libraries, for example the ‘rIntChron’ package gives direct access to IntChron-indexed databases in R (CITE).

- **p3k14c** (CITE) instead compiles multiple source databases into a single flat file dataset, with a similar level of coverage to c14bazAAR. The major advantage of this approach is that the data is made internally consistent and has been manually cleaned to an extent, which makes it particularly well-suited to global analyses. The downside is that without the continuous link to the source databases present in the c14bazAAR and IntChron, it can only be kept up to date manually with periodic re-releases. An accompanying package (CITE<empty citation>) provides direct access to the p3k14c dataset in R.

XRONOS blends aspects of all three approaches to achieve the same aim of providing access to the global radiocarbon data through a common interface. Like c14bazAAR and IntChron, it is a ‘metadatabase’ that draws from existing data resources and maintains an explicit link to them. But like p3k14c, it integrates these into a single database and applies data curation processes to harmonise them and improve the quality of the information. It has a wider scope than c14bazAAR or IntChron, as it mirrors rather than directly retrieves the source data (allowing us to use resources that aren’t openly published), and does not rely on the authors of these sources to implement a common specification. It also goes beyond the functionality of p3k14c by providing systems for the continuous ingestion and curation of new data as it is published. However in general we see the approaches as complementary. A c14bazAAR parser for XRONOS is in development (<https://github.com/ropensci/c14bazAAR/pull/150>), and we also aim to provide an IntChron interface to XRONOS’ data in the near future. New data and corrections from p3k14c are incorporated into XRONOS as they are released.

XRONOS is our answer to Kintigh’s call (Kintigh 2006) for digital infrastructures that don’t just provide access to chronological data but enables researchers to “archive, access, integrate, and mine disparate data sets”. It parallels and draws inspiration from several similar initiatives in archaeology, such as And outside of it:

Beyond radiocarbon

- Dendro
- Typology
- Other radiometric methods
- Application-specific
- XRONOS aims to bring all these together, starting with radiocarbon, dendro-, and typology

Concept

XRONOS inherits its basic structure from RADON (CITE<empty citation>), with a database-backed web application and a data model that separates radiocarbon dates, contextual information, and sites. Our overall aims in developing XRONOS is to bring this model, which RADON has operated on for XX years, up to date, to generalise it to other types of chronometric information, and to transform it from an online database to a data infrastructure that supports the continuous ingestion, curation, and open dissemination of archaeological chronologies from diverse sources.

Design goals

Data model

- FIGURE: conceptual scheme

At the base of the XRONOS data model are sets of spatiotemporal coordinates or, as we call them, *chrons*. In an archaeological context, we conceptualise a chron as an assertion linking human activity with a particular point in space and time. Our data model currently encompasses three types of chron: radiocarbon dates, typological dates (e.g. ‘Early Neolithic’) and dendrochronological dates. However we anticipate that the concept will accommodate other types of absolute and relative dating techniques, as the scope of the database expands.

Chrons are conceptually useful because they emphasise that different types of archaeological ‘dates’, drawn from different sources, have essentially the same information content: the location of an event in space and time. We thereby avoid privileging certain sources of chronological data over (as might be the case if, for example, we treated ‘period’ as a fixed attribute of a site) and can accommodate contradictory (e.g. differences of opinion on typological classification). This is important given that XRONOS aspires to be an authoritative ‘backbone’ with a global scope, so we cannot realistically impose a single chronological scheme or resolve conflicting information provided by specialists. They are useful practically because they expose a common interface for attributes that all types of chronological information share, such as a *terminus post quem* (TPQ), *terminus ante quem* (TAQ), and midpoint estimate. This allows applications that use XRONOS’ data model (including XRONOS itself) to collate chronological data from multiple sources, without necessarily having to be aware of the peculiarities of each type of dating.

In order to unify chronological information in the form of a chron, we need a common chronological ‘coordinate system’. The natural choice is a *calendar probability distribution*, which expresses the probability that an event occurred as a function of time on a calendric scale. Most archaeologists are familiar with working with this kind of representation in the form of calibrated radiocarbon dates, but it can be extended and generalised to essentially any kind of chronological information (CITE). For example, in aoristic analysis (CITE), a periodic time estimate (e.g. the event occurred in the Neolithic) is conceptualised as a uniform probability distribution over the timespan between the known start and end dates of that period. A similar model is used in OxCal (OxCal) to integrate prior chronological information from diverse sources. In practical terms, this model means that the canonical representation the time component of any chron in XRONOS’, regardless of source, is a probably distribution over the set of calendar years (arbitrarily measured in years Before Present) in which it could have plausibly occurred. Further statistics, e.g. a midpoint estimate or TPQ/TAQ range, can be derived from this distribution using well-known methods. In this way, we can support many different types of date and much of the implementation of XRONOS can be agnostic to the source of chronological information.

Chrons are located in space through association to a *sample* – the physical object from which a chronological determination was made. The location of samples is represented with geographical coordinates and an associated coordinate reference system (CRS), though since in practice the precise location of single samples is rarely available, this property is usually inherited from the site. We also record relevant metadata on the nature of the sample. For radiocarbon dates, for example, we follow established conventions (Millard 2014) in recording the type (e.g. charcoal, charred seed) and, where applicable, taxonomic designation (e.g. TODO, TODO)

of the organic material used for dating. For typological dates, an ideal scenario would be for the sample to represent the particular object from which an inference was made (e.g. ‘Natufian’ might be inferred from ‘lunate-type microlith’). In practice, the best we can glean from most published datasets is the type of material used (e.g. ‘pottery’, ‘lithics’). The same sample can be associated with multiple chrons, including different types of chron. This is useful, for example, for representing replicate radiocarbon dates on the same sample, or radiocarbon dates and dendrochronological made on the same section of wood for wiggle-matching.

Further contextual information is associated with *contexts* and *sites*. The site is the primary geographic container for chronological information. As already mentioned, we typically record the spatial location of chrons using this entity, though it is possible to modify this by providing specific coordinates at the sample level. Sites also have attributes describing their conventional name or names in different languages and are associated with a flexible ‘site type’ typology that combines information on their form and function.

A context represents the specific find-context of a sample, e.g. an architectural feature, stratigraphic unit, or phase. Since the units and conventions for recording such information vary greatly between different regions and archaeological traditions—and XRONOS is designed with global data in mind—we leave the question of what a context precisely represents open, and only record an unstandardised, free text label for it. Crucially, however, contexts can have a self-referential association to other contexts belonging to the same site. This allows it to encode arbitrary relational structures between contexts, whether they be hierarchical (e.g. phases and sub-phases) or graphical (e.g. stratigraphic). In this way, it can serve as a foundation for chronological modelling.

The series of relations [chron] > sample > context > site links the chronological and contextual sides of the XRONOS data model. Each step is a many-to-many association, meaning for example that it is possible to attach multiple chrons to the same sample (e.g. replicated radiocarbon dates on the same material), multiple *types* of chrons to the same sample (e.g. radiocarbon dates on tree-rings for wiggle-matching). Since this kind of information is rarely systematically recorded in our source databases, there are currently few actual records that make use of this feature of the data model. However, we hope it will provide a foundation for ... ?

Metadata is incorporated into XRONOS’ data model at the level of the individual records (e.g. all records store their data of creation and last modification) and through two additional nodes: bibliographic references and versions. Bibliographic references store information on the source of a record in the standard BibTeX (CITE<empty citation>) format and can be linked through many-to-many associations to sites or chrons. Versions are a special type of record that are associated with all other records (including bibliographic references) and store the previous versions of those records as a series of changesets. In this way all changes to data are recorded and can be reconstructed (or reversed) precisely. This ‘paper trail’ also stores contextual metadata, e.g. who made the change and why. It also means that records that are deleted can be reviewed or restored from their stored version history, which is never discarded. Together these two systems provide a transparent record where the data in XRONOS comes from and how it has been altered, which we view as essential in a scientific data infrastructure.

Linked data

The XRONOS data model presents several opportunities to link to other resources as linked open data (LOD) (CITE). We use controlled vocabularies in the form of the GBIF Backbone Taxonomy (for taxonomic descriptions of samples) via its API, and in future plan to extend this to for example site types (using the Getty and/or DAI). Standardising field using these vocabularies to automatically then link the two entities. Beyond this, there is also a concordance between XRONOS entities and other specialised data infrastructures in archaeology, such as Perio.do (CITE<empty citation>, mapping to typological chrons) and gazetters like Pleiades or Vici (CITE<empty citation>, mapping to site names). Moving beyond archaeology, Wikidata (CITE<empty citation>) already includes many of the concepts represented in the XRONOS model (e.g. archaeological sites, Q). Linking to Wikidata is especially useful because it disseminates—and thereby preserves—the data compiled in XRONOS beyond the narrow field of archaeology. It also allows us to enrich the database with contextual information that otherwise be beyond our scope and resources, for example embedding multilingual descriptions of sites from Wikipedia articles on site records where the site has been linked to a Wikidata item.

Conversely, we encourage others to use XRONOS as a linked open data resource by providing stable URLs and a machine-readable interface for every record. Such usages could look like, for example, using a XRONOS URL as a canonical representation of a single radiocarbon date (e.g.). We also plan to implement an IntChron ((CITE)) interface to XRONOS, to allow it to be indexed through that existing standard.

Implementation

Following a short pilot project in 2019, the first phase of development of XRONOS was completed in 2021–2024, supported by a project grant from the Swiss National Science Foundation. The web interface (<https://xronos.ch>) has been publicly accessible since July 2021. Though we envisage XRONOS as a continuously-developing open source project and ‘living database’, the following offers a snapshot of progress at the end of our first grant-funded implementation phase. We do not aim to be comprehensive, but rather to describe some key elements of XRONOS’ current implementation that illustrate how our concept has been realised in practice.

Software architecture

The XRONOS data model is implemented as a relational database using the free and open source database management system PostgreSQL. However, apart from backups and other routine maintenance procedures, all interaction with the database is via a web application, which thus forms the core of XRONOS’ architecture. The XRONOS web application is written in Ruby and uses CRUD (Create, Read, Update, Delete) and MVC (Model, View, Controller) patterns as implemented in the Ruby on Rails framework. The choice of this somewhat traditional or ‘boring’ architecture—as opposed to the non-relational/semantic technologies that are more *à la mode* in the digital humanities (Fan 2018; Hyvönen 2020; Schloen and Prosser 2023)—was purposeful, motivated by a sense that it offers a better chance of longevity and maintainability. Our aim is for the software architecture to be as boring as possible so that the scientific contents can be as interesting as possible.

The XRONOS web application exposes two distinct user interfaces: a graphical user interface accessed through a web browser; and an application programming interface (API). Both interfaces follow a REST (Representational State Trans-

fer) pattern (Verborgh et al. 2015), where each resource (e.g. a single radiocarbon date, a single bibliographic reference, or a single user) is statelessly mapped to a single address. Users can then interact with resources at these addresses using a predictable and uniform interface based on HTTP verbs. For example, the radiocarbon date RTD-8904 is represented by the address <https://xronos.ch/c14s/156205>. Users can view information on this resource by sending a GET request to that address, regardless of which interface they are using, and authorised users can modify it using POST, PATCH, or DELETE requests. The bibliographic reference associated with this date (Richter et al. 2017) is similarly represented at the address <https://xronos.ch/references/17778>, and can be accessed at that address using the same interface as the radiocarbon date. These uniform REST interfaces are another example of a boring architectural choice that make it easier for us to enrich the scientific contents of XRONOS, by adding new types of modular resource that represent new scientific entities.

This basic REST pattern is augmented by seven ‘actions’ (following the standard pattern in Rails application) that express different ways of interacting with a resource: index, show, destroy, new, create, edit, and update. The ‘show’ action represents interaction with a single resource, as described above. The ‘index’ action, which lists resources of a given type (e.g. <https://xronos.ch/c14s> for radiocarbon dates), is worth special mention because it is through this that the filtering logic at the core of XRONOS’ two interfaces is implemented. By passing a query as HTTP GET parameters to the index action of a resource, the list returned the user is modified to only include records that match that query. For example, https://xronos.ch/sites?site%5Bcountry_code%5D=CH (the part of the URL after the ? character encodes the SQL WHERE clause `country = 'CH'` as a GET parameter) lists sites in Switzerland. More complex queries can be executed using nested parameters. For example, <https://xronos.ch/c14s?c14%5Bsample%5D%5Bmaterial%5D%5Bname%5D=charcoal> (encoding that the `c14` table should be joined to the `material` table via `sample`, followed by the WHERE clause `material.name = 'charcoal'`) lists radiocarbon dates obtained from charcoal samples. Uniquely, index actions can also respond with the result in a tabular data format (i.e. .csv).

Data ingestion and curation

The chronological data in XRONOS comes from a variety of sources, including published structured datasets in repositories and journal supplements, other on-line databases, literature review, and direct input from collaborators. Our aim is not just to ‘mirror’ these sources as they are, but integrate them into a single curated and continuously updated database. For the purposes of ingestion, we classify data resources into three categories: static resources, such as supplementary data in published papers, which are imported once; versioned resources, updated on a periodic basis, which we import after each new version; and live resources, which are continuously updated and therefore continuously imported. Records of each type are imported into XRONOS in as close to their original state as possible, i.e. without any corrections or standardisation applied. This ensures that any subsequent changes (even immediate, automatic ones) are entered into the record’s version history, so that the source of any deviations or potential errors can always be reconstructed. The version history also records the direct source of the data for attribution purposes. In addition, a bibliographic reference to the original resource is attached to ensure that the source is clearly attributed even if the record is merged with another one.

Once ingested, we apply a number of automated and semi-automated quality control processes to integrate new data into the existing database. Controlled vocabularies are used in a number of places in the data model (FIG<empty citation>?), and we use thesauruses to automatically standardise these fields as much as possible. For example, the taxonomic description of samples is controlled using GBIF's backbone taxonomy (CITE<empty citation>), and we also use a thesaurus service provided by GBIF to automatically change variant or obsolete taxonomic names to the canonical version. If the system is not able to standardise a field using the available thesaurus, it is flagged for manual correction. A wide variety of other potential data quality issues (e.g. missing data on what country a site is in) are also flagged for human review by this system, which can often be semi-automated (e.g. suggesting close matches in the thesaurus or the country indicated by the record's coordinates).

A final critical component of XRONOS' data curation system is duplicate handling. We import data from many overlapping resources (many of which incorporate each other either in whole or in part), so duplicate records are common. The end result of standardising and correcting a record is also often to create a duplicate: e.g. the same sample imported from one source as 'oak' but another as '*Quercus* sp.' will become a duplicate pair as '*Quercus*', and thus be recognised as a single sample. Such exact duplicates can be merged automatically, with the oldest record becoming the authoritative version, but detecting fuzzier duplicated information (e.g. differences in the spelling of site names) has proved a more difficult problem. As of writing there are therefore still many duplicate records in XRONOS that need to be manually resolved, but we hope to automate much more of this work in the future.

User interfaces

The graphical user interface (GUI) to XRONOS, accessed through a web browser (e.g. at <https://xronos.ch>), uses REST resources and actions as the building blocks for various interfaces through which users can browse, search, retrieve, and analyse chronometric data. Each action on each resource is represented by a page, though not all of these are publicly accessible. Pages representing REST resources directly are supplemented by a number of synthetic interfaces, for example the 'data browser' (<https://xronos.ch/data>), which facilitates more complex filtering, or the search interface (<https://xronos.ch/search>). The GUI also includes several resources which are not part of XRONOS' scientific data model, for example documentation pages, user profiles and news articles; these are currently not exposed in the API.

Access to various 'backstage' interfaces for creating, editing, and deleting data, and monitoring data quality is managed using a user permissions system. Currently only authorised users affiliated with the XRONOS project can access these, but in the future we intend to support open registration and expose editing interfaces to all authenticated users. For this reason, there is no sharp division between a 'public' and 'private' areas – viewing/querying data and editing/curating data share the same architecture and interface patterns.

The XRONOS API uses the same addresses as the web-based GUI (with the exception of some of the synthetic interfaces mentioned above) but responds with machine-readable data in JSON format, rather than a HTML page. This response can be triggered by appending `.json` to the address or by including a HTTP content-format header in the request. Though users can make such requests manually and parse the data with one of several off-the-shelf tools, the primary

intended uses of this interface is to provide access for 1) programmatic clients to XRONOS and 2) other web services. The XRONOS R package (CITE) is an example of a programmatic client; it uses the API to facilitate direct querying and retrieval of data from XRONOS in the R statistical programming language (CITE), which is widely used for computational applications in archaeology (CITE). Similar libraries could be developed in other programming environments used for scientific computing, such as Python or Haskell. The API also provides the foundation for other web services to access XRONOS directly, to embed chronological information in other contexts or otherwise make use of its data resources.

An overarching principle of this software architecture is that all interaction with XRONOS' data store, and as much of the data processing and 'business logic' of responding to REST requests as possible, is directed through the same server-side routines. First and foremost, this allows us to provide multiple interfaces (i.e. the GUI and API, perhaps more in the future) without duplicating these elements of our codebase. It also improves accessibility for users accessing XRONOS through devices with limited processing capability or through text-only browsers. More broadly, avoiding reliance on client-side processing, e.g. with Javascript or Web Assembly (WASM)—which would be the other option—allows us keep our client interfaces simple (in most cases plain, semantic HTML pages and self-contained stylesheets) and therefore, we hope, sustainable in the face of constantly-evolving client-side technologies and standards. It does have the weakness that, in practical terms, XRONOS is difficult to run and relies on the continued existence of a maintained external server. We have however tried to mitigate this by providing clearly-documented source code and regular data dumps so that, if our instance of XRONOS disappears, or if one simply does not want to use it, it is possible for others to host a XRONOS server of their own. As the sustainability of scientific software and data infrastructures is a pressing problem (a point we will return to in the conclusion), in the future it may be desirable to support further decentralisation through, for example, a federated server-server model (CITE).

Evaluation

- XRONOS is useable, use XRONOS!
- Immediate development goals
- The sustainability challenge

References

- Fan, Lai-Tze. 2018. "On the Value of Narratives in a Reflexive Digital Humanities." *Digital Studies / Le champ numérique* 8, no. 1 (1 2018). <https://doi.org/10.16995/dscn.285>.
- Hyvönen, Eero. 2020. "Using the Semantic Web in Digital Humanities: Shift from Data Publishing to Data-Analysis and Serendipitous Knowledge Discovery." *Semantic Web* 11, no. 1 (January 1, 2020): 187–193. <https://doi.org/10.3233/SW-190386>.
- Kintigh, Keith. 2006. "The Promise and Challenge of Archaeological Data Integration." Publisher: Cambridge University Press, *Am. Antiq.* 71, no. 3 (July): 567–578. Accessed August 15, 2019. <https://doi.org/10.1017/S0002731600039810>. <https://www.cambridge.org/core/journals/american-antiquity/article/promise-and-challenge-of-archaeological-data-integration/037763619B121990D59B585523826A03>.
- Millard, Andrew R. 2014. "Conventions for Reporting Radiocarbon Determinations." Publisher: Cambridge University Press, *Radiocarbon* 56 (2): 555–559. Accessed February 2, 2018. <https://doi.org/10.2458/56.17455>. <https://www.cambridge.org/core/journals/radiocarbon/article/conventions-for-reporting-radiocarbon-determinations/E4077EC6F1EE1C90C5170309E2C0CF9B>.

- Richter, Tobias, Amaia Arranz-Otaegui, Lisa Yeomans, and Elisabetta Boaretto. 2017. "High Resolution AMS Dates from Shubayqa 1, Northeast Jordan Reveal Complex Origins of Late Epipalaeolithic Natufian in the Levant." *Sci. Rep.* 7, no. 1 (December 5, 2017): 17025. <https://doi.org/10.1038/s41598-017-17096-5>.
- Schloen, Sandra R., and Miller C. Prosser. 2023. "The Case for a Database Approach." In *Database Computing for Scholarly Research: Case Studies Using the Online Cultural and Historical Research Environment*, edited by Sandra R. Schloen and Miller C. Prosser, 25–73. Cham: Springer International Publishing. https://doi.org/10.1007/978-3-031-46696-0_2.
- Verborgh, Ruben, Seth van Hooland, Aaron Straup Cope, Sebastian Chan, Erik Mannens, and Rik Van de Walle. 2015. "The Fallacy of the Multi-API Culture: Conceptual and Practical Benefits of Representational State Transfer (REST)." *Journal of Documentation* 71, no. 2 (January 1, 2015): 233–252. <https://doi.org/10.1108/JD-07-2013-0098>.